



[< Return to Classroom](#)

NLP on Financial Statements

REVIEW

CODE REVIEW

HISTORY

Meets Specifications

Congratulations!

You have demonstrated an excellent understanding of the concepts in Project 5 and the ability to implement those concepts in Python.

- It has been a pleasure reviewing your project.

Great News!

- you are more than 1/2 way through this nanodegree!!
- Good luck on your last 3 projects!

Here are some additional resources which can help you take your newly acquired skills even further!

Blueshift (Formerly Quantopian) is a site for learning and practicing Quantitative Python Programming that provides Python workspace similar to what you've been using in your nanodegree.

- [Blueshift Tutorials](#) : 3 full tutorials by Quantopian to assist with quant workflow.

10-Ks

The function `get_documents` extracts the documents from the text.

Well done!

- you set up a nice start pattern and a nice end pattern for using regular expressions to pull just the documents from an input string.

```
1 import re
2
3
4 def get_documents(text):
5     """
6     Extract the documents from the text
7
8     Parameters
9     -----
10    text : str
11        The text with the document strings inside
12
13    Returns
14    -----
15    extracted_docs : list of str
16        The document strings found in `text`
17    """
18
19    # TODO: Implement
20    extracted_docs = []
21    for start_idx, end_idx in zip([x.end() for x in re.compile(r'<DOCUMENT>').finditer(text)], [x.start() for x
22    extracted_docs.append(text[start_idx:end_idx])
23    return extracted_docs
24
25
26 project_tests.test_get_documents(get_documents)
```

Tests Passed

The function `get_document_type` returns the document type lowercased.

Document Type is being determined properly.

- You have properly constructed a regular expression for obtaining the document type.
- You are returning the document type as a lowercase string.

```
1 def get_document_type(doc):
2     """
3     Return the document type lowercased
4
5     Parameters
6     -----
7     doc : str
8         The document string
9
10    Returns
11    -----
12    doc_type : str
13        The document type lowercased
14    """
15
16    # TODO: Implement
17    return [x[len('<TYPE>'):] for x in re.compile(r'<TYPE>[^\n]+').findall(doc)][0].lower()
18
19
20 project_tests.test_get_document_type(get_document_type)
```

```
20 project_tests.test_get_document_type(get_document_type)
```

Tests Passed

Preprocess the Data

The function `lemmatize_words` lemmatizes verbs.

Verbs are being properly lemmatized!

- You are properly using **WordNetLemmatizer** to lemmatize verbs in the list of words.
- Nice use of a list comprehension!

```
1 from nltk.stem import WordNetLemmatizer
2 from nltk.corpus import wordnet
3
4
5 def lemmatize_words(words):
6     """
7     Lemmatize words
8
9     Parameters
10    -----
11    words : list of str
12           List of words
13
14    Returns
15    -----
16    lemmatized_words : list of str
17                       List of lemmatized words
18    """
19
20    # TODO: Implement
21
22    return [WordNetLemmatizer().lemmatize(word, pos='v') for word in words]
23
24
25 project_tests.test_lemmatize_words(lemmatize_words)
```

Tests Passed

Analysis on 10ks

The function `get_bag_of_words` generates a bag of words from documents.

Sentiment Bag of Words is properly created.

- You have nicely used **CountVectorizer** on the incoming `sentiment_words` and then used that to transform the incoming docs, creating a “bag of words”.

```

1 from collections import defaultdict, Counter
2 from sklearn.feature_extraction.text import CountVectorizer
3
4
5 def get_bag_of_words(sentiment_words, docs):
6     """
7     Generate a bag of words from documents for a certain sentiment
8
9     Parameters
10    -----
11    sentiment_words: Pandas Series
12        Words that signify a certain sentiment
13    docs : list of str
14        List of documents used to generate bag of words
15
16    Returns
17    -----
18    bag_of_words : 2-d Numpy Narray of int
19        Bag of words sentiment for each document
20        The first dimension is the document.
21        The second dimension is the word.
22    """
23
24    # TODO: Implement
25    bag_of_words = CountVectorizer(vocabulary=sentiment_words).fit_transform(docs).toarray()
26    return bag_of_words
27
28
29 project_tests.test_get_bag_of_words(get_bag_of_words)

```

Tests Passed

The function `get_jaccard_similarity` calculates the jaccard similarities for neighboring documents.

You are correctly calculating the Jaccard Similarity on the Bag of Words

- Nice use of list comprehension and zip to calculate Jaccard similarities for neighboring documents!

```

1 from sklearn.metrics import jaccard_similarity_score
2
3
4 def get_jaccard_similarity(bag_of_words_matrix):
5     """
6     Get jaccard similarities for neighboring documents
7
8     Parameters
9     -----
10    bag_of_words : 2-d Numpy Narray of int
11        Bag of words sentiment for each document
12        The first dimension is the document.
13        The second dimension is the word.
14
15    Returns
16    -----
17    jaccard_similarities : list of float
18        Jaccard similarities for neighboring documents
19    """
20
21    # TODO: Implement
22    bag_of_words = bag_of_words_matrix.astype(bool)
23    jaccard_similarities = [jaccard_similarity_score(a, b) for a, b in zip(bag_of_words, bag_of_words[1:])]
24    return jaccard_similarities
25
26
27 project_tests.test_get_jaccard_similarity(get_jaccard_similarity)

```

Tests Passed

The function `tfidf` generate TFIDF vectors for each document.

Yes!

– You used `TfidfVectorizer` (in a similar manner to how you used `CountVectorizer` for your `bag_of_words`) to generate the sentiment TFIDF from the 10-k documents using the sentiment words as the terms.

```

1 from sklearn.feature_extraction.text import TfidfVectorizer
2
3
4 def get_tfidf(sentiment_words, docs):
5     """
6     Generate TFIDF values from documents for a certain sentiment
7
8     Parameters
9     -----
10    sentiment_words: Pandas Series
11        Words that signify a certain sentiment
12    docs : list of str
13        List of documents used to generate bag of words
14
15    Returns
16    -----
17    tfidf : 2-d Numpy Narray of float
18        TFIDF sentiment for each document
19        The first dimension is the document.
20        The second dimension is the word.
21    """
22
23    # TODO: Implement
24    tfidf = TfidfVectorizer(vocabulary = sentiment_words).fit_transform(docs).toarray()
25    return tfidf
26
27
28 project_tests.test_get_tfidf(get_tfidf)

```

Tests Passed

The function `get_cosine_similarity` calculates the cosine similarities for each neighboring TFIDF vector/document.

Cosine Similarity is Accurately calculated using the TFIDF values!

```

1 from sklearn.metrics.pairwise import cosine_similarity
2
3
4 def get_cosine_similarity(tfidf_matrix):
5     """
6     Get cosine similarities for each neighboring TFIDF vector/document
7
8     Parameters
9     -----
10    tfidf : 2-d Numpy Narray of float
11        TFIDF sentiment for each document
12        The first dimension is the document.
13        The second dimension is the word.
14
15    Returns
16    -----
17    cosine_similarities : list of float
18        Cosine similarities for neighboring documents
19    """

```

```
20  
21 # TODO: Implement  
22 return cosine_similarity(tfidf_matrix[0:], tfidf_matrix[1:])[0].tolist()  
23  
24  
25 project_tests.test_get_cosine_similarity(get_cosine_similarity)
```

Tests Passed

In case you are interested . . .

Here is a solution that uses numpy's diag function and does not require looping:

```
return list(np.diag(cosine_similarity(tfidf_matrix, tfidf_matrix), k=1))
```

 [DOWNLOAD PROJECT](#)

[RETURN TO PATH](#)