

[◀ Return to Classroom](#)

Generate Faces

REVIEW

CODE REVIEW

HISTORY

Meets Specifications

Kudos ! I think you've done a perfect job of implementing a Deep Convolutional GAN to Generate Faces. It's very clear that you have a good understanding of the basics.

Further Reading: One of the biggest problem GAN researchers face (you yourself must have experienced) with standard loss function is, the quality of generated images do not correlate with loss of either G or D. Since both the network are competing against each other, the losses fluctuate *a lot*. This was solved in early 2017 with introduction of [Wasserstein GANs](#). Do read it up.

Finally, have a look at [this amazing library](#) by Google for training and evaluating Generative Adversarial Networks.

And Congratulations 🎉 once again, as you've completed this project for the Deep Learning Nanodegree. I hope you had great experience with Udacity and will continue taking other Nanodegree too.

Required Files and Tests

The project submission contains the project notebook, called "d1nd_face_generation.ipynb".

The `iPythonNB` and helper files are included.

All the unit tests in project have passed.

Great work. Unit testing is one of the most reliable methods to ensure that your code is free from all bugs without getting confused with the interactions with all the other code. If you are interested, you can [read up more](#) and I hope that you will continue to use unit testing in every module that you write to keep it clean and speed up your development.

But always keep in mind, that unit tests cannot catch every issue in the code. So your code could have bugs even though unit tests pass.

Build the Neural Network

The function `model_inputs` is implemented correctly.

Correct.

Placeholders is the building block in computation graph of any neural net (especially in **tensorflow**).

Often I find students confused between `tf.Variable` and `tf.placeholder`. [This answer](#) gives correct usecase for both.

The function `discriminator` is implemented correctly.

Overall you did a fine job implementing the `Discriminator` as a simple convolution network.

Let me illustrate the pros of the architecture you chose.

Pros

- `tf.variable_scope('discriminator', reuse=reuse)` was essential to this part for two reasons. Firstly, to make sure all the variable names start with `discriminator`. This will help out later when training the separate networks. Secondly, the discriminator will need to share variables between the fake and real input images using `reuse`.
- You chose not to use pooling layers to decrease the spatial size. Max pooling generates sparse gradients, which affects the stability of GAN training. We generally use Average Pooling or Conv2d + stride.
- Correctly used Leaky ReLU. As explained above we never want sparse gradients (~ 0 gradients). Therefore, we use a leaky ReLU to allow gradients to flow backwards through the layer unimpeded.
- Used Batch normalization. We initialize the BatchNorm Parameters to transform the input to zero mean/unit variance distributions but as the training proceeds it can learn to transform to `x` mean and `y` variance, which might be better for the network. [This post](#) is an awesome read to understand BatchNorm to it's core.
- Using dropout in discriminator so that it is less prone to learning the data distribution.
- Using a sigmoid for output layer.

Tips

- Use custom weight initialization. Xavier init is proposed to work best when working with GANs.

The function generator is implemented correctly.

Most of the suggestions are same for both Generator and Discriminator.

Let me (again) illustrate the pros of the architecture you chose.

Pros

- `Tanh` as the last layer of the generator output. This means that we'll have to normalize the input images to be between -1 and 1.

Tips

- Try decreasing the width of layers from 512 -> 64. In context of GANs, a sharp decline in number of filters for Generator helps produce better results.

The function `model_loss` is implemented correctly.

Perfect.

Now that was the trickiest part (and my personal favorite in GAN :)

The function `model_opt` is implemented correctly.

Clean and concise.

Neural Network Training

The function `train` is implemented correctly.

- It should build the model using `model_inputs`, `model_loss`, and `model_opt`.
- It should show output of the `generator` using the `show_generator_output` function

Good work normalizing inputs using `x *= 2`.

You might want to watch up this talk on [How to train a GAN](#) by one of the author of original DCGAN paper and corresponding [write-up](#).

The parameters are set reasonable numbers.

Given your network architecture, the choice of hyper-parameter are reasonable.

Tips

- You selected a good value for `beta1`. Here's a [good post](#) explaining the importance of beta values and which value might be empirically better. Also try lowering it even further, ~0.1 might even produce better results.
- An important point to note is, batch size and learning rate are linked. If the batch size is too small then the gradients will become more unstable and would need to reduce the learning rate.

The project generates realistic faces. It should be obvious that images generated look like faces.

Woah ! Now that is something 🖼️+1:

This simple DCGAN model is generating faces too close to real human faces. Wonder what will happen when trained on for 5-10 epochs ?

Fun Read: Ever heard about the analogy “king – man + woman = queen”? [Read this up](#) on how to perform the same analogies on images, using simple arithmetic in latent space.

 [DOWNLOAD PROJECT](#)

RETURN TO PATH

Rate this project

START